

Machine Learning Project: Predicting Wine Quality

1 Introduction

This project aims to predict wine quality based on objective physicochemical characteristics, such as acidity, density, pH, and residual sugar and alcohol content [1]. To achieve this, we will develop a machine learning model trained on a dataset of red and white variants of the Portuguese “Vinho Verde” wine. Each wine gets a quality score from 1 to 10, 10 being the best [1].

The rest of this report is structured as follows: Section 2 formalizes the application into a machine learning problem; Section 3 describes data preprocessing, correlation analysis, splitting strategy, and the Ridge Regression model; Section 4 documents the use of AI tools; and Section 5 presents references and the repository appendix.

2 Problem Formulation

We formulate this task as a supervised machine learning regression problem. Given an array of physicochemical measurements for a wine sample, our goal is to predict its integer quality score $y_i \in \mathbb{Z}$. Each data point $x_i \in \mathbb{R}^{12}$ represents a single wine sample defined by 12 feature variables:

#	Feature Name	Data Type	Description [1]
1	fixed acidity	Continuous	Non-volatile acids (g/dm ³)
2	volatile acidity	Continuous	Acetic acid content (g/dm ³)
3	citric acid	Continuous	Freshness/flower acid content (g/dm ³)
4	residual sugar	Continuous	Remaining sugar post-fermentation (g/dm ³)
5	chlorides	Continuous	Salt concentration (g/dm ³)
6	free sulfur dioxide	Continuous	Free SO ₂ preventing microbial growth (mg/dm ³)
7	total sulfur dioxide	Continuous	Total free and bound SO ₂ (mg/dm ³)
8	density	Continuous	Liquid density (g/cm ³)
9	pH	Continuous	Acidic/basic scale (0–14)
10	sulphates	Continuous	SO ₂ gas additive (g/dm ³)
11	alcohol	Continuous	Alcohol content (% by volume)
12	is white	Binary	Variant indicator (1 = white, 0 = red)

3 Methods

3.1 Preparing the data

The dataset consists of $N = 6497$ total wine samples, split between red wine ($N_{\text{red}} = 1599$) and white wine ($N_{\text{white}} = 4898$). We concatenated both subsets into a single master dataset and introduced an engineered binary feature **is white** ($x_{i,12} \in \{0,1\}$) to allow the linear model to account for chemical baseline differences between wine types increasing our feature vector dimension to 12.

Figure 1 shows the sample count split between red and white wines. Figure 2 illustrates the distribution of target quality scores y ; most ratings cluster around scores 5 and 6, showing that extreme scores are relatively rare, meaning that the data is not heavily skewed in either direction. This helps in making an accurate model. In addition it is worth mentioning that no wine received a score of 0, 1, 2 or 10. These values are, therefore, omitted from the histogram.

3.2 Data correlation

Inspection of the Pearson correlation matrix (Figure 3) reveals notable collinearity among several feature pairs:

- **Sulfur Dioxide Metrics:** **free sulfur dioxide** and **total sulfur dioxide** exhibit a strong positive correlation ($r = 0.72$), as free SO_2 is a direct subset of total SO_2 . We considered engineering a ratio feature $\text{sulfur}_{\text{ratio}} = \frac{\text{free_sulfur_dioxide}}{\text{total_sulfur_dioxide}}$, which may be tested in Stage 2.
- **Density Relationships:** **density** correlates negatively with **alcohol** ($r = -0.69$) and positively with **residual sugar** ($r = 0.55$), reflecting physical fluid properties.
- **Wine Variant Correlations:** The **is white** feature strongly correlates with **total sulfur dioxide** ($r = 0.70$) and negatively with **volatile acidity** ($r = -0.65$).

We retain all 12 features in our Stage 1 feature vector $X \in \mathbb{R}^{6497 \times 12}$. Retaining the full feature set preserves maximum physical variance and avoids potential underfitting caused by aggressive feature deletion.

3.3 Data splitting

We split the dataset into three disjoint partitions:

- **Training Set (60%, $N_{\text{train}} = 3898$):** Used for fitting parameter weights.
- **Validation Set (20%, $N_{\text{val}} = 1299$):** Reserved for hyperparameter tuning (λ) and model selection.
- **Test Set (20%, $N_{\text{test}} = 1300$):** Retained untouched for final generalization error evaluation.

This provides a large enough training set without sacrificing validation or test capacity. To prevent feature scale disparities from dominating the L_2 penalty, we apply z-score standardization:

$$\tilde{x}_{ij} = \frac{x_{ij} - \mu_j}{\sigma_j}$$

To avoid data leakage, feature means (μ_j) and standard deviations (σ_j) are calculated strictly on the training set and subsequently applied to transform the validation and test sets.

3.4 Model Selection and Loss Function

To model continuous wine quality while handling correlated features without manual feature dropping, we select **Ridge Regression** as our model.

3.4.1 Hypothesis Space The hypothesis space consists of linear predictors mapping the 12-dimensional feature vector x_i to a predicted score $\hat{y}_i$:

$$h(x_i, w) = w_0 + \sum_{j=1}^{12} w_j x_{ij} = w_0 + w^T x_i$$

where w_0 is the intercept/bias term and $w = (w_1, \dots, w_{12})^T$ is the weight vector.

3.4.2 Loss Function and Regularization We optimize the model using regularized Mean Squared Error (MSE) with an L_2 -norm penalty:

$$L(w) = \frac{1}{2n} \sum_{i=1}^n (h(x_i, w) - y_i)^2 + \frac{\lambda}{2} \|w\|_2^2$$

where $\|w\|_2^2 = \sum_{j=1}^{12} w_j^2$ penalizes large weight magnitudes, and $\lambda \geq 0$ is the regularization hyperparameter. The L_2 penalty stabilizes weight estimation in the presence of collinearity (e.g., between sulfur dioxide features), mitigating overfitting and improving generalization performance without losing accuracy or data.

4 Use of AI

During this stage AI (Google Gemini) was used for the following tasks:

- **Brainstorming** Exploring different datasets and their suitability for our purposes.
- **LaTeX and Markdown formatting** Helping while writing our report to make sure that the document looks correct and does not have any syntax errors.
- **Report Structure and Grammar** Reviewing the grammar of the the report and its structure and suggesting possible fixes.

All of the code, data-analysis and report writing was done by hand by both members of the team. All of the AI suggestions were considered and reviewed carefully together.

5 Appendices and References

5.1 Visualization




Figure 1: Red vs White Wine Distribution




Figure 2: Wine Quality Histogram




Figure 3: Pearson Correlation Matrix

5.2 Source Code

The full, runnable, source code is available in a public, anonymous GitHub repository: <https://github.com/Computerplayer4/wine-machine-learning>

For reading convenience, it is also available below [2][3][4][5][6]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Reading the data
url_red =
    "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv"
url_white =
```

```

    "https://archive.ics.uci.edu/ml/machine-learning-databases/wine-quality/winequality-white.

# Load data into dataframes
df_red = pd.read_csv(url_red, sep=';')
df_white = pd.read_csv(url_white, sep=';')

# Add the `is white` binary feature such that 0 = Red and 1 = White
df_red['is white'] = 0
df_white['is white'] = 1

# Combine the two dataframes into one ignoring
# the indexing as it's not useful to preserve it
df_wine = pd.concat([df_red, df_white], axis=0, ignore_index=True)

# Form the feature matrix X and label vector y
X = df_wine.drop(columns=['quality'])
y = df_wine['quality'].to_numpy()

# Splitting the dataset into training (60 %), validation (20 %) and test (20 %) sets
X_train, X_val_test, y_train, y_val_test =
    train_test_split(X, y, test_size=0.4, random_state=42)
X_val, X_test, y_val, y_test =
    train_test_split(X_val_test, y_val_test, test_size=0.5, random_state=42)

# Normalize the data to zero mean and unit variance
# We only use training data when calculating mu and sigma as
# we want to prevent data leakage
# These values will be used with training and validation data only
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)
X_test_scaled = scaler.transform(X_test)

# Generate EDA statistics
print("--- DATASET SUMMARY ---")
print(f"Total samples (N): {len(df_wine)}")
print(f"Total white wine samples: {len(df_white)}")
print(f"Total red wine samples: {len(df_red)}")
print(f"Feature vector dimension (d): {X.shape[1]}")
print(f"Train size: {len(X_train)}, Val size: {len(X_val)}, Test size: {len(X_test)}\n")

print(f"--- FEATURE SUMMARY STATISTICS ---")
print(df_wine.describe().T[['mean', 'std', 'min', 'max']])

# Generate wine counts figure

wine_counts = df_wine['is white'].value_counts()
plt.figure(figsize=(6, 4))

```

```

# Here 0 is Red and 1 is White as before
plt.bar(['Red Wine', 'White Wine'], [wine_counts.get(0, 0),
    wine_counts.get(1, 0)], color=['#800020', '#F0E68C'], edgecolor='black')
plt.title('Dataset split: Red vs. White Wine Samples')
plt.xlabel('Wine Type')
plt.ylabel('Number of Samples (Count)')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.savefig('../assets/wine_type_split.png')
plt.show()

# Generate score histogram
plt.figure(figsize=(8, 5))
plt.hist(df_wine['quality'], bins=range(0,11), align='left', rwidth=0.85,
    color='purple', edgecolor='black', alpha=0.8)
plt.title('Distribution of Wine Quality Scores')
plt.xlabel('Quality Score (Rating)')
plt.ylabel('Frequency (Number of Samples)')
plt.xticks(range(0, 10))
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.savefig('../assets/quality_histogram.png')
plt.show()

# Generate Pearson Correlation Matrix Heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(df_wine.corr(), annot=True, fmt=".2f", cmap="coolwarm")
plt.title("Pearson Correlation Matrix - Combined Wine Quality")
plt.tight_layout()
plt.savefig("../assets/correlation_matrix.png", dpi=300)
plt.show()

```

5.2 References

- [1] P. Cortez, A. Cerdeira, F. Almeida, T. Matos, and J. Reis. “Wine Quality,” UCI Machine Learning Repository, 2009. [Online]. Available: <https://doi.org/10.24432/C56S3T>.
- [2] “pandas documentation — pandas 3.0.5 documentation.” <https://pandas.pydata.org/docs/>
- [3] “NumPy documentation — NumPy v2.5 Manual.” <https://numpy.org/doc/stable/>
- [4] “seaborn: statistical data visualization — seaborn 0.13.2 documentation.” <https://seaborn.pydata.org/>
- [5] “Using Matplotlib — Matplotlib 3.11.2 documentation.” <https://matplotlib.org/stable/users/index>
- [6] “User Guide,” Scikit-learn. https://scikit-learn.org/stable/user_guide.html